

# Software Management and Containers in HPC

**Pablo Mata Aroco - BU-ISCIII**

**29-04 de octubre de 2025**

**1ª edición**



# Overview

---

1. Software management basics.
2. Software managers: EasyBuild, Conda-Micromamba
3. Introduction to Virtualization. Containers: Docker & Singularity

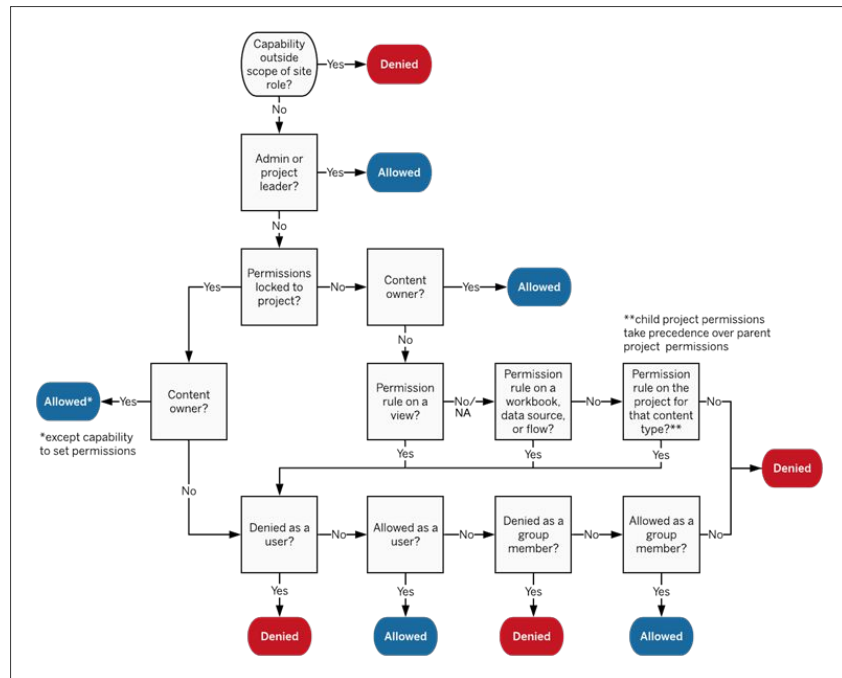
# Overview

1. Software management basics
2. Software managers: EasyBuild, Conda-Micromamba
3. Introduction to Virtualization. Containers: Docker & Singularity

# 1. Software management basics: Permissions

**Permissions are rules that control what a user can do on a computer system. They determine:**


- Who can read (view a file or program).
- Who can write (modify or delete).
- Who can execute (run a program).



# 1. Software management basics: Permissions

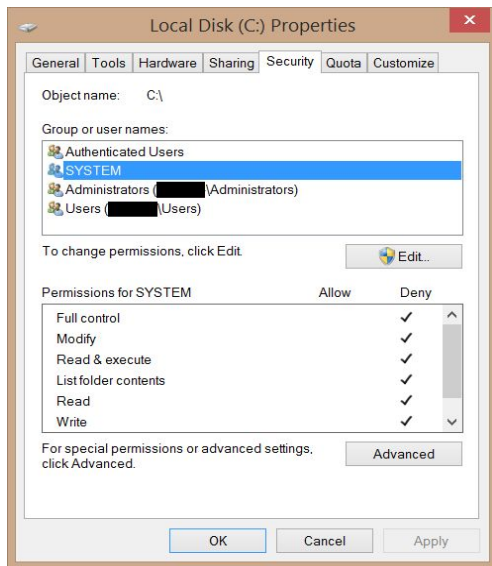
They exist to **protect** the operating system, applications, and data from accidental damage or malicious use:

- **Security:** Prevent malware or mistakes from harming other users.
- **Stability:** Prevent one user's bad installation from breaking shared applications.
- **Consistency:** Admins control which versions of software are installed cluster-wide.
- **Performance:** HPC software often needs to be built with optimized compilers and libraries.

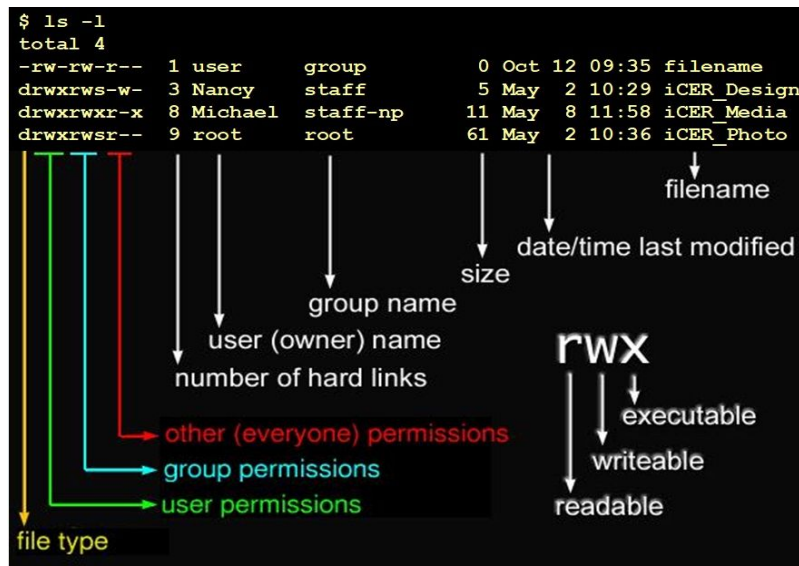


# 1. Software management basics: Windows vs Linux

## Windows



## Linux

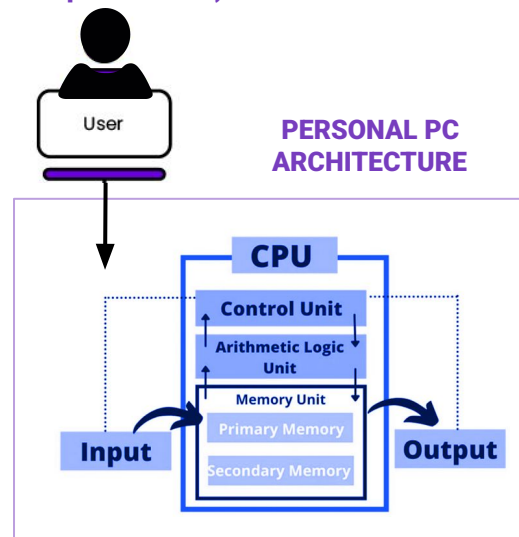


# 1. Software management: HPC vs Personal PC

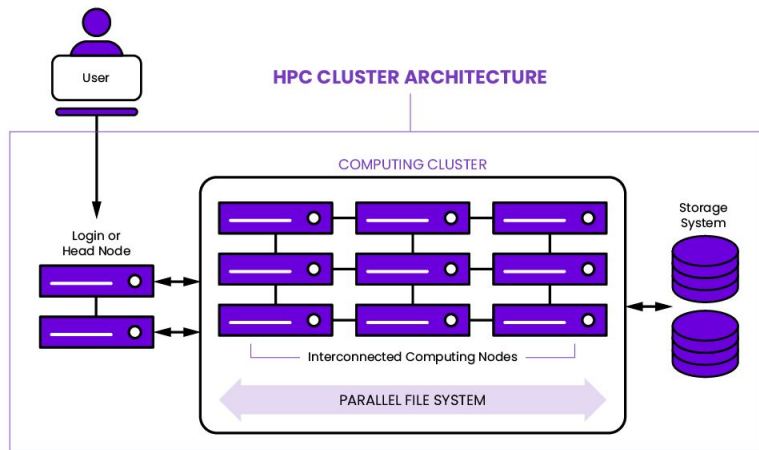
On a **personal computer** (for example, your Windows laptop) installing software is straightforward:

- You open a browser, download an **installer** (.exe or .msi), run it, and the program installs locally.
- **You usually have administrator rights** (even if you don't notice), so you can change almost any system configuration.
- **You work on a single machine**, with a fixed operating system and hardware (CPU, RAM, graphics card...).

(Admin permissions)



# 1. Software management: HPC vs Personal PC



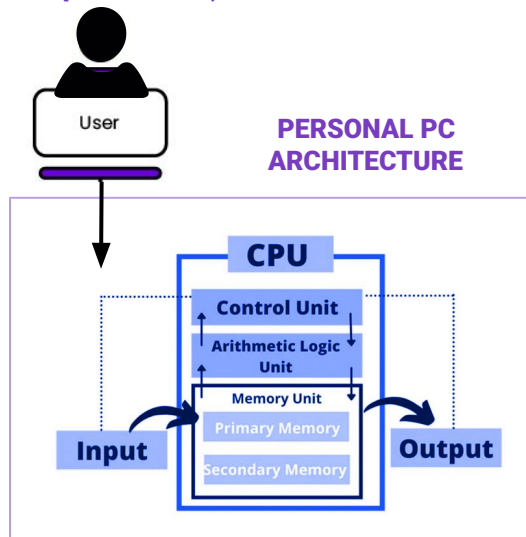
On a High-Performance Computing (HPC) system, the situation is very different:

- It's a **multi-user environment**
- The hardware consists of compute nodes that **cannot be accessed by normal users**
- You have **no root permissions**: you cannot install software system-wide, change OS settings...
- The goal is to optimize for the cluster hardware and ensure reproducibility, not just “get a program running”

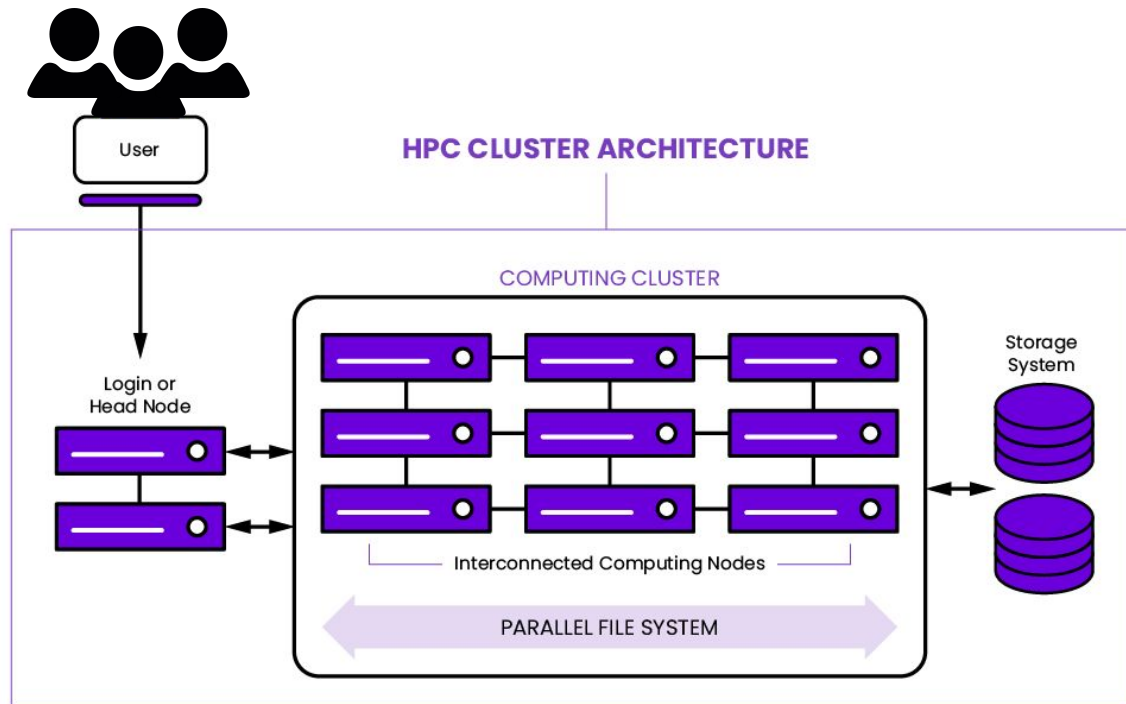


# 1. Software management: HPC vs Personal PC

(Admin permissions)



**HPC CLUSTER ARCHITECTURE**



# 1. Software management: HPC principles

On a personal machine, you can afford to “break” your setup and reinstall. On an HPC the situation is completely different:

- Overwriting system libraries could **break unrelated software for all the other users.**
- You must be able to **reproduce exactly an analysis months or years later** with the same configuration.
- In case of **unrestricted installation** a malicious or careless user could **install malware.**

That's why HPC systems:

- **Organize software into modules, virtual environments or containers** so you can choose versions without interfering with others.
- This latter approach also helps to **keep reproducibility.**
- **Contained privileges:** Users are sometimes free to install packages but only in their home or scratch directories, where they can't affect other users.

# 1. Software management: HPC principles

## 1. Security Reasons

An HPC is a shared multi-user environment — dozens, hundreds, or even thousands of people run jobs on the same nodes. If everyone could install or modify system software freely, it would be a security nightmare.

Key security risks if installation were unrestricted:

- **Malware** **injection:**  
A malicious or careless user could install software that steals credentials, mines cryptocurrency, or damages data.
- **Privilege** **escalation:**  
If someone installs a program with a known vulnerability, attackers could exploit it to gain root access.
- **Tampering** **tools:**  
If `ssh`, `ls`, or `python` were replaced with compromised versions, all users would be at risk.
- **Inadvertent** **breakage:**  
Without bad intent, a user might overwrite a library with an incompatible version, causing other users' jobs to crash.

Example: Imagine installing a Python library that silently sends computation results to an external server. On your laptop, you'd only risk your own data — on an HPC, you could leak other researchers' sensitive datasets.

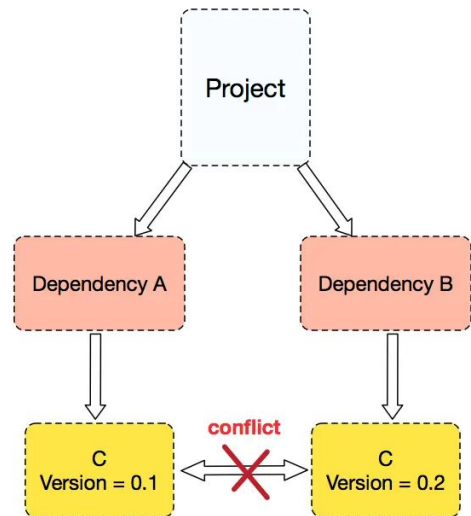
# 1. Software management: HPC principles

## 2. Maintainability Reasons

HPC admins have to maintain stability across hundreds of nodes, often for years.

Challenges if users installed software freely:


- Version** **chaos:**  
 If every user installed different versions of `gcc`, `cuda`, or `numpy` in system directories, no one could guarantee reproducibility.
- Dependency** **conflicts:**  
 Overwriting system libraries could break unrelated software



# 1. Software management: HPC principles

How admins handle it instead:

- **Centralized software management:** HPC teams use tools like EasyBuild or Spack to build and deploy software in a controlled way across all nodes.
- **Environment modules:** Users load software via `module load python/3.10` instead of installing it themselves.
- **Containers:** Singularity or Docker (where allowed) lets users run custom environments without touching the base system.



# HPC vs Personal PC: Summary

Feature	Personal PC (Windows)	HPC
<b>Simultaneous users</b>	Typically one or a few, separate sessions	Dozens or hundreds of active users at the same time.
<b>Permissions</b>	User can install and modify system software	No root; install only in your home space if needed
<b>Software installation</b>	Run graphical installers. No risk of conflict with other users.	Use modules, environment managers, or compile with tools like EasyBuild, Conda, Mamba
<b>Hardware</b>	Single CPU (multi-core), limited RAM	Multiple nodes, each with many cores and large memory; high-speed interconnect
<b>Purpose</b>	General use (office, games, etc.)	Large-scale data processing or computationally intensive tasks
<b>Program execution</b>	Run directly on your machine	Submit jobs to a scheduler/queue; executed on compute nodes
<b>Updates</b>	Automatic, controlled by the user	Planned by administrators to avoid disrupting jobs
<b>Version management</b>	Install or uninstall at will	Multiple versions coexist. You choose which to use

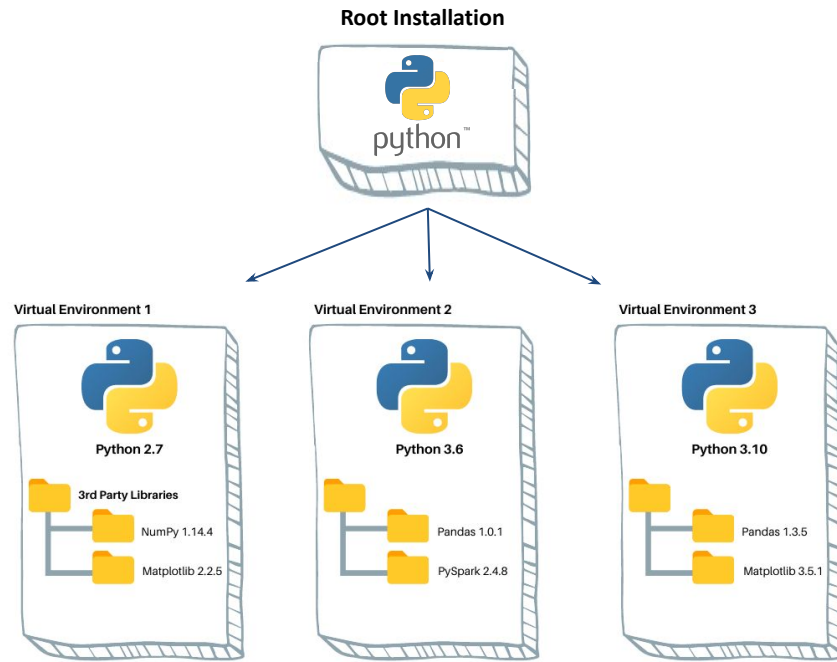
# Overview

1. Software management basics.
2. Software managers: EasyBuild, Conda-Mamba
3. Introduction to Virtualization. Containers: Docker & Singularity

### 3. Introduction to Virtual Environments

#### Virtual Environments.

- A **virtual environment** is an **isolated workspace** on your computer where you can install software **without interfering** with the global system.
- **Examples:** Python `venv`, Conda, Micromamba environments, R virtual libraries.
- **Pro:** No root permissions needed.
- **Con:** Doesn't virtualize the OS, system or hardware. Only manages the software layer.



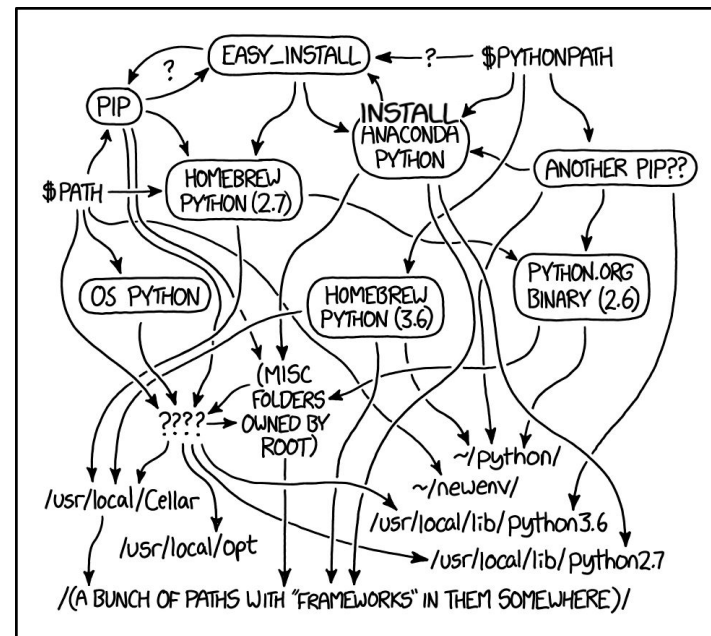


# 3. Introduction: Virtual Environments

## Virtual Envs usage in HPC

On an **HPC cluster**, you usually **don't have admin/root permissions**. That means:

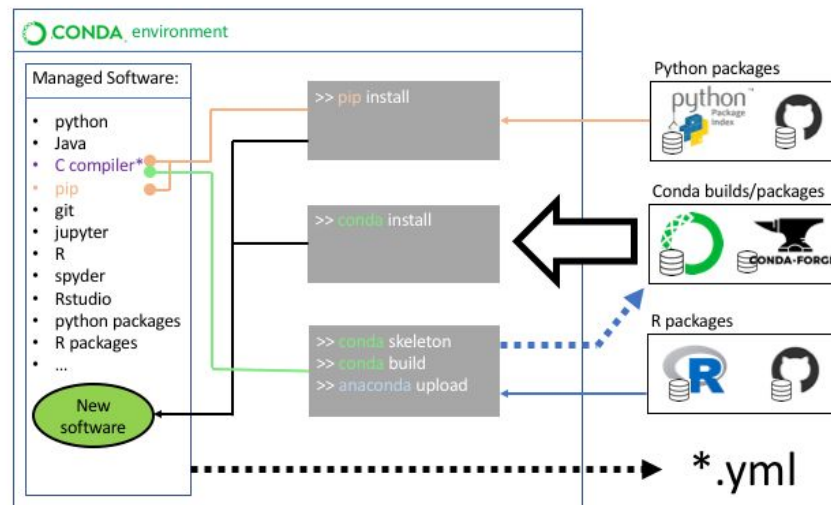
- You **cannot install system-wide software** (like `pip` `install` globally).
- Users shares the same filesystem**: installing packages globally would cause **conflicts**.
- Even inside a virtual environment its difficult to avoid conflicts. **Software managers** exist for the sole reason to **simplify this process**.



# 3. Software managers: Conda/Mamba

## 3. Virtual Environments (Software-Level)

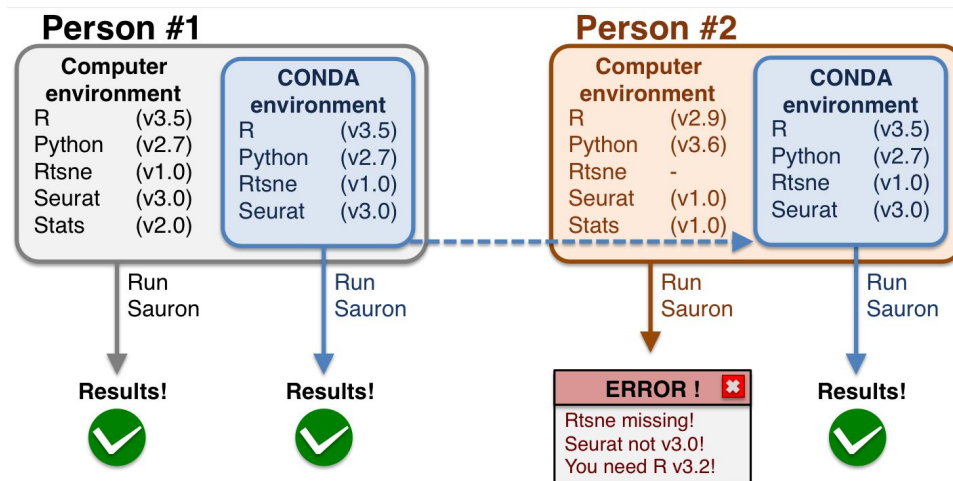
- **Pro:** No root permissions needed, isolates packages and dependencies **inside the same OS** without affecting system installations.
- **Con:** Doesn't virtualize the OS, system or hardware. Only manages the software layer.
- **HPC usage:** Extremely common for running different Python/R environments on shared clusters without interfering with system-wide libraries.
- **Example:** Python `venv`, Conda, Micromamba environments, R virtual libraries.



### 3. Software managers: Conda

#### Conda

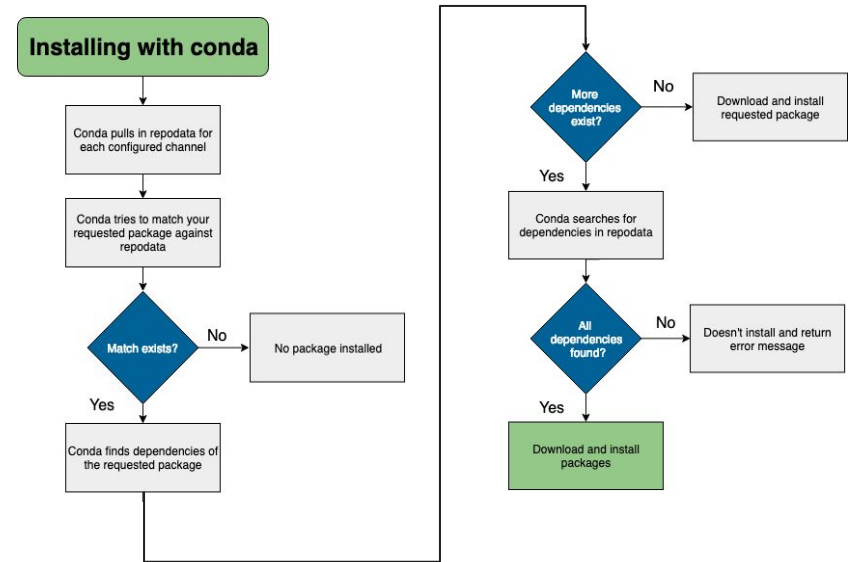

- **Conda** is a **package/environment manager**. Originally created for Python but works with **any language** (R, C/C++, Fortran...).
- Creates **isolated environments**, each with its own Python/Language version and libraries.
- Manages **dependencies**: ensures all required libraries are installed in compatible versions.



# 3. Software managers: Conda

## Install with conda. Channels

- Not everything can be installed with *conda install*. The desired software must be accessed through a **conda channel**
- Conda channels** are **repositories** (URLs) where Conda looks for packages.
- The base channel is *defaults*, but there are community-driven channels like *conda-force* or *bioconda*. Anyone can create a channel.



### 3. Software managers: Mamba

#### Mamba

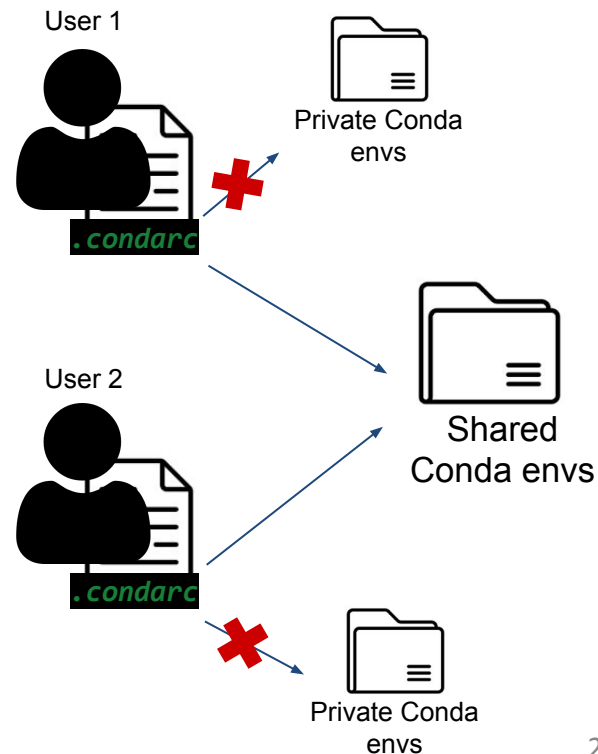
- A **drop-in replacement for Conda** written in C++ instead of Python.
- Uses the same package repositories and commands as Conda.
- Designed to solve Conda's main weakness: **slowness**, specially in dependency resolution.
- **Micromamba**: A lightweight (10mb), and self-contained version of mamba (installed in our HPC).



### 3. HPC shared environments

#### HPC shared virtual envs

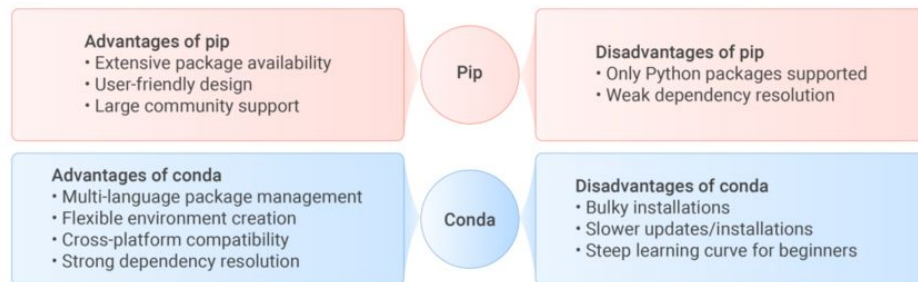
- If each user creates its own environments in the HPC, there may be **differences in results due to software versions**.
- Using shared environments **avoids differences in analysis** results due to environment differences.
- In our HPC, **all users can access the same *modules***, for softwares like **Nextflow** or **Singularity**
- You can also **change conda's configuration file (*.condarc*)** to use the same list of environments as other users.



### 3. Other package managers: Pip

#### Pip

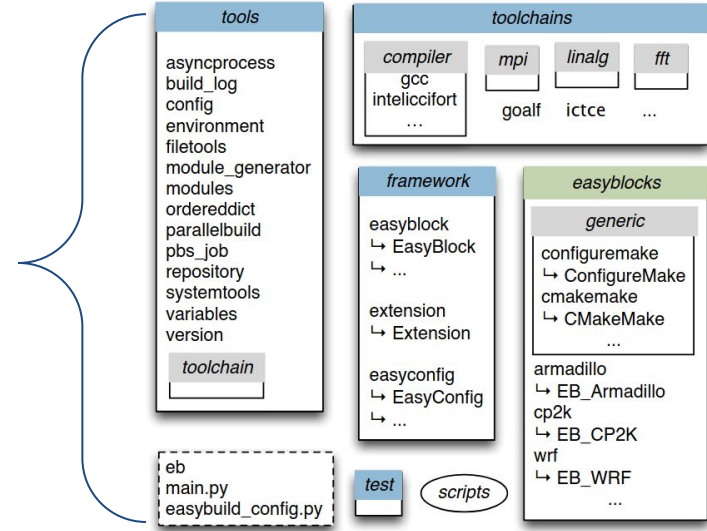

- Pip is the **default package manager for Python**. It is the **official tool** for installing packages from the Python Package Index (<https://pypi.org/>), which holds a **vast catalogue of packages**.
- Pip's sole purpose is to **install, update, and remove Python packages from PyPI**. Nonetheless, it can install binary packages (wheels) or software from source.
- Pip itself **does not have built-in environment management** like conda, but you can install it in a conda environment.



### 3. Software managers: EasyBuild

#### Easybuild.

- EasyBuild is an **open-source framework** designed specifically for **building and installing scientific software on HPC systems**.
- Our HPC uses this framework to manage certain global software for all users (e.g. Nextflow or singularity)
- **Encapsulates software in modules**, which can be loaded via: `module load <module-name>`






# Overview

1. Software management basics.
2. Software managers: EasyBuild, Conda-Mamba
3. Introduction to Virtualization. Containers: Docker & Singularity

### 3. Introduction to virtualization

#### What is Virtualization?

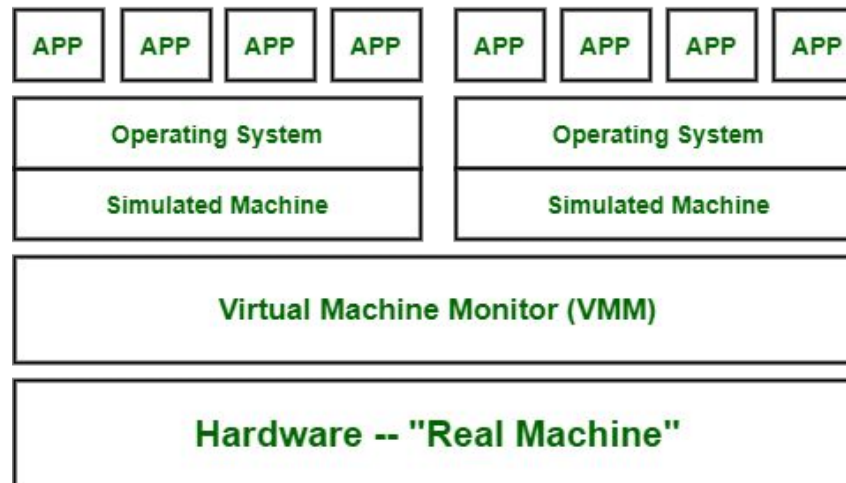
- Virtualization is the process of creating a **virtual version of a computing resource**.
- It **behaves like the real thing, but is actually a software** on top of physical hardware.
- Instead of interacting directly with the physical server, **you interact with a virtual replica** that is **isolated from your system**.



# 3. Introduction to virtualization: Virtual Machines

## A. Full Virtual Machines

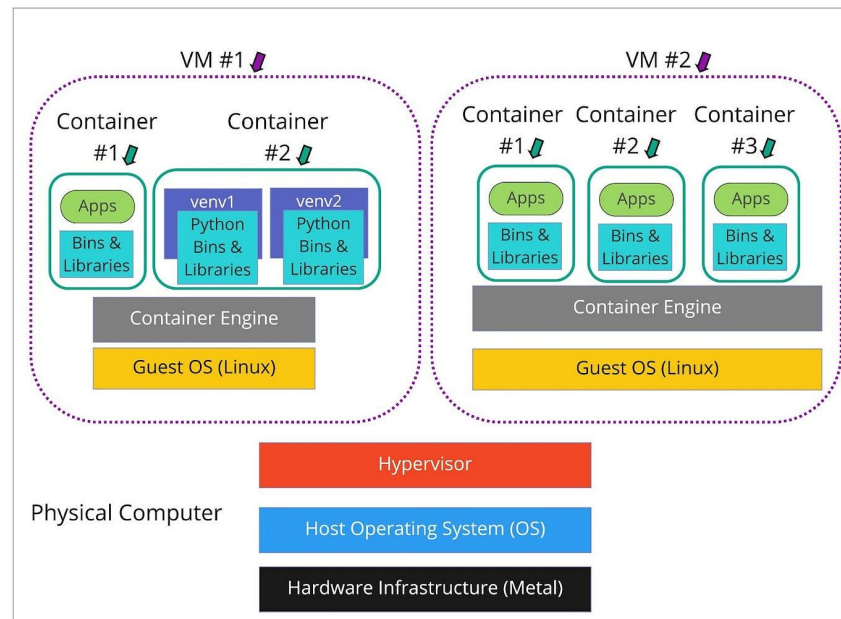
- A **Virtual Machine (VM)** is a software-based emulation of a physical computer, including OS and hardware components using a **hypervisor**.
- **Pro:** Total isolation. They **can run a completely different operating system**.
- **Con:** **Very slow** due to full hardware emulation.
- **HPC usage:** **Rare**, mainly for testing environments or isolated, security-sensitive workloads.
- Example: Running **Ubuntu in VirtualBox** on a Windows laptop.



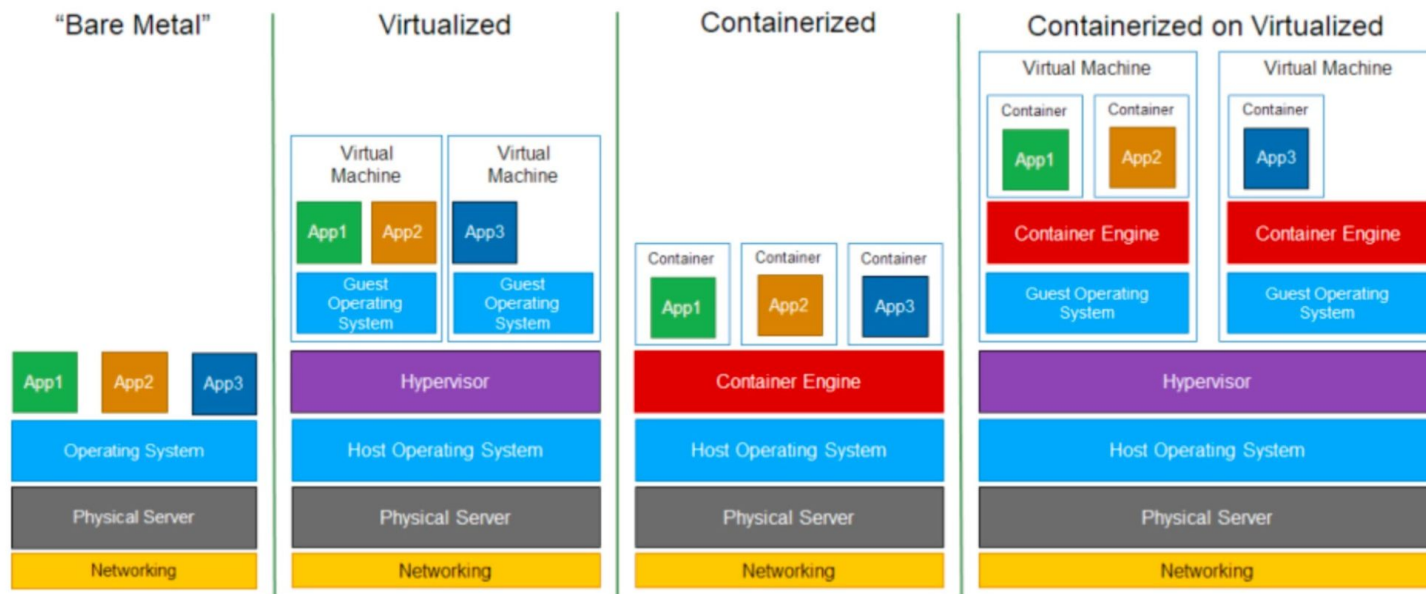
# 3. Introduction to virtualization: Containers

## B. OS-level Virtualization (Containers)

- A container is a lightweight, **standalone package** that **bundles together an application and everything it needs to run**, but **shares the host system's kernel** instead of emulating hardware.
- **Pro:** Near-native performance, lightweight.
- **Con:** Shares the host OS kernel, **cannot simulate a different OS than the host's**.
- **HPC usage:** **Very common** for packaging software to run without dependency issues.
- Examples: Docker, Singularity.



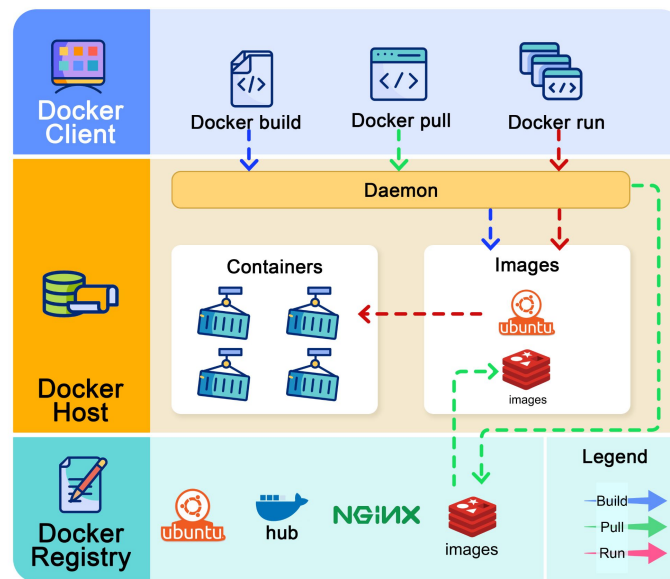
### 3. Introduction to virtualization: VMs vs Containers



## 3.1 Containers: Docker & Singularity

### Docker

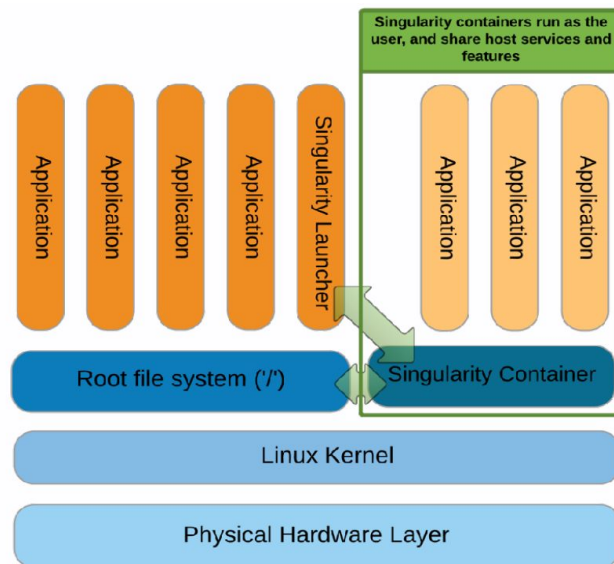
- Docker is a popular platform designed for building, sharing, and running applications in **containers**.
- By default, Docker containers run as the **root user** on the host system. This can pose a **security risk in multi-user environments** like an HPC.
- Unfortunately, we cannot directly use Docker in our HPC.



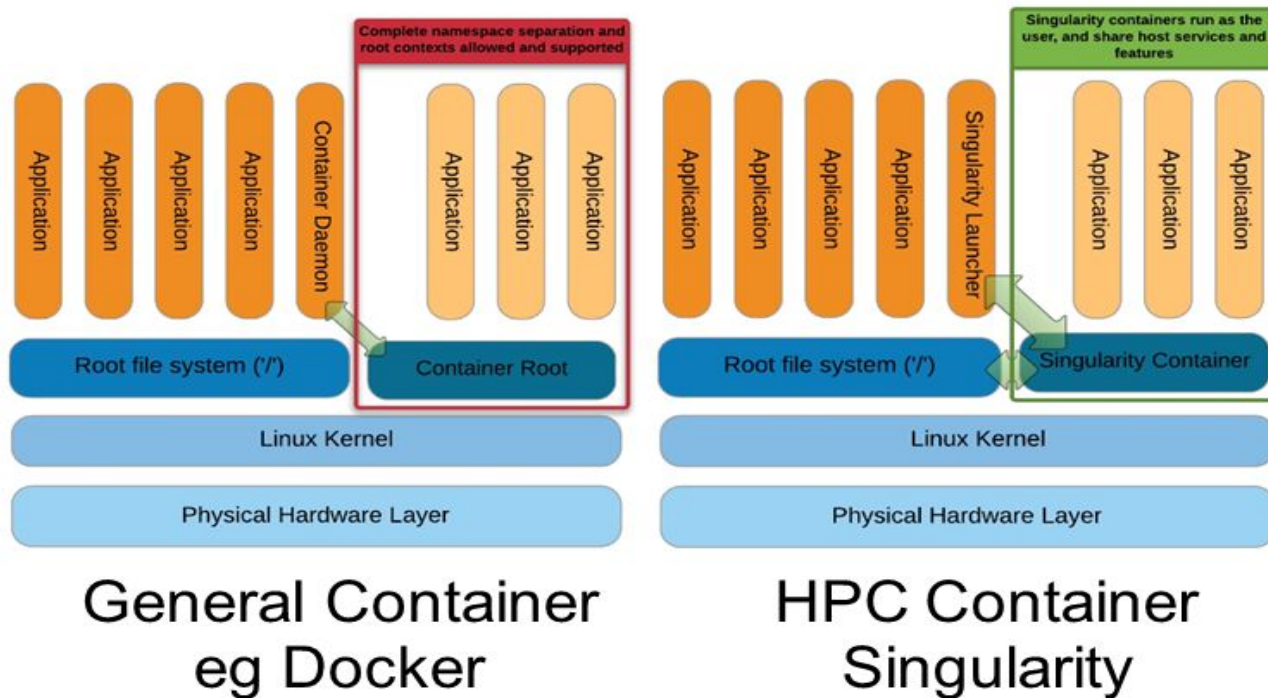
## 3.1 Containers: Docker & Singularity

### Singularity

- Singularity is a container platform specifically **designed for high-performance computing (HPC)** and scientific research environments.
- Singularity containers **run as the same user that launched them** on the host system. This is a **security feature for HPC clusters**.
- It's our go-to to create and run containers in the HPC
- Galaxy's <https://depot.galaxyproject.org/singularity> includes a lot of singularity **containers for bioinformatic analysis**:



## 3.1 Containers: Docker & Singularity





# Essential takeaways

- You cannot install anything on the HPC that needs admin permissions.
- Prioritize **modules over virtual environments** (More efficient)
- If you only need to **execute one certain task** (e.g. run FastQC) use a **singularity container**
- If you need to **interact with files or data** dynamically use **virtual environments** (Micromamba/Pip)
- **DON'T RUN HEAVY JOBS IN THE LOGIN NODE!!**



# Practical exercises

**The best way to understand everything is with some hands-on exercises.**

